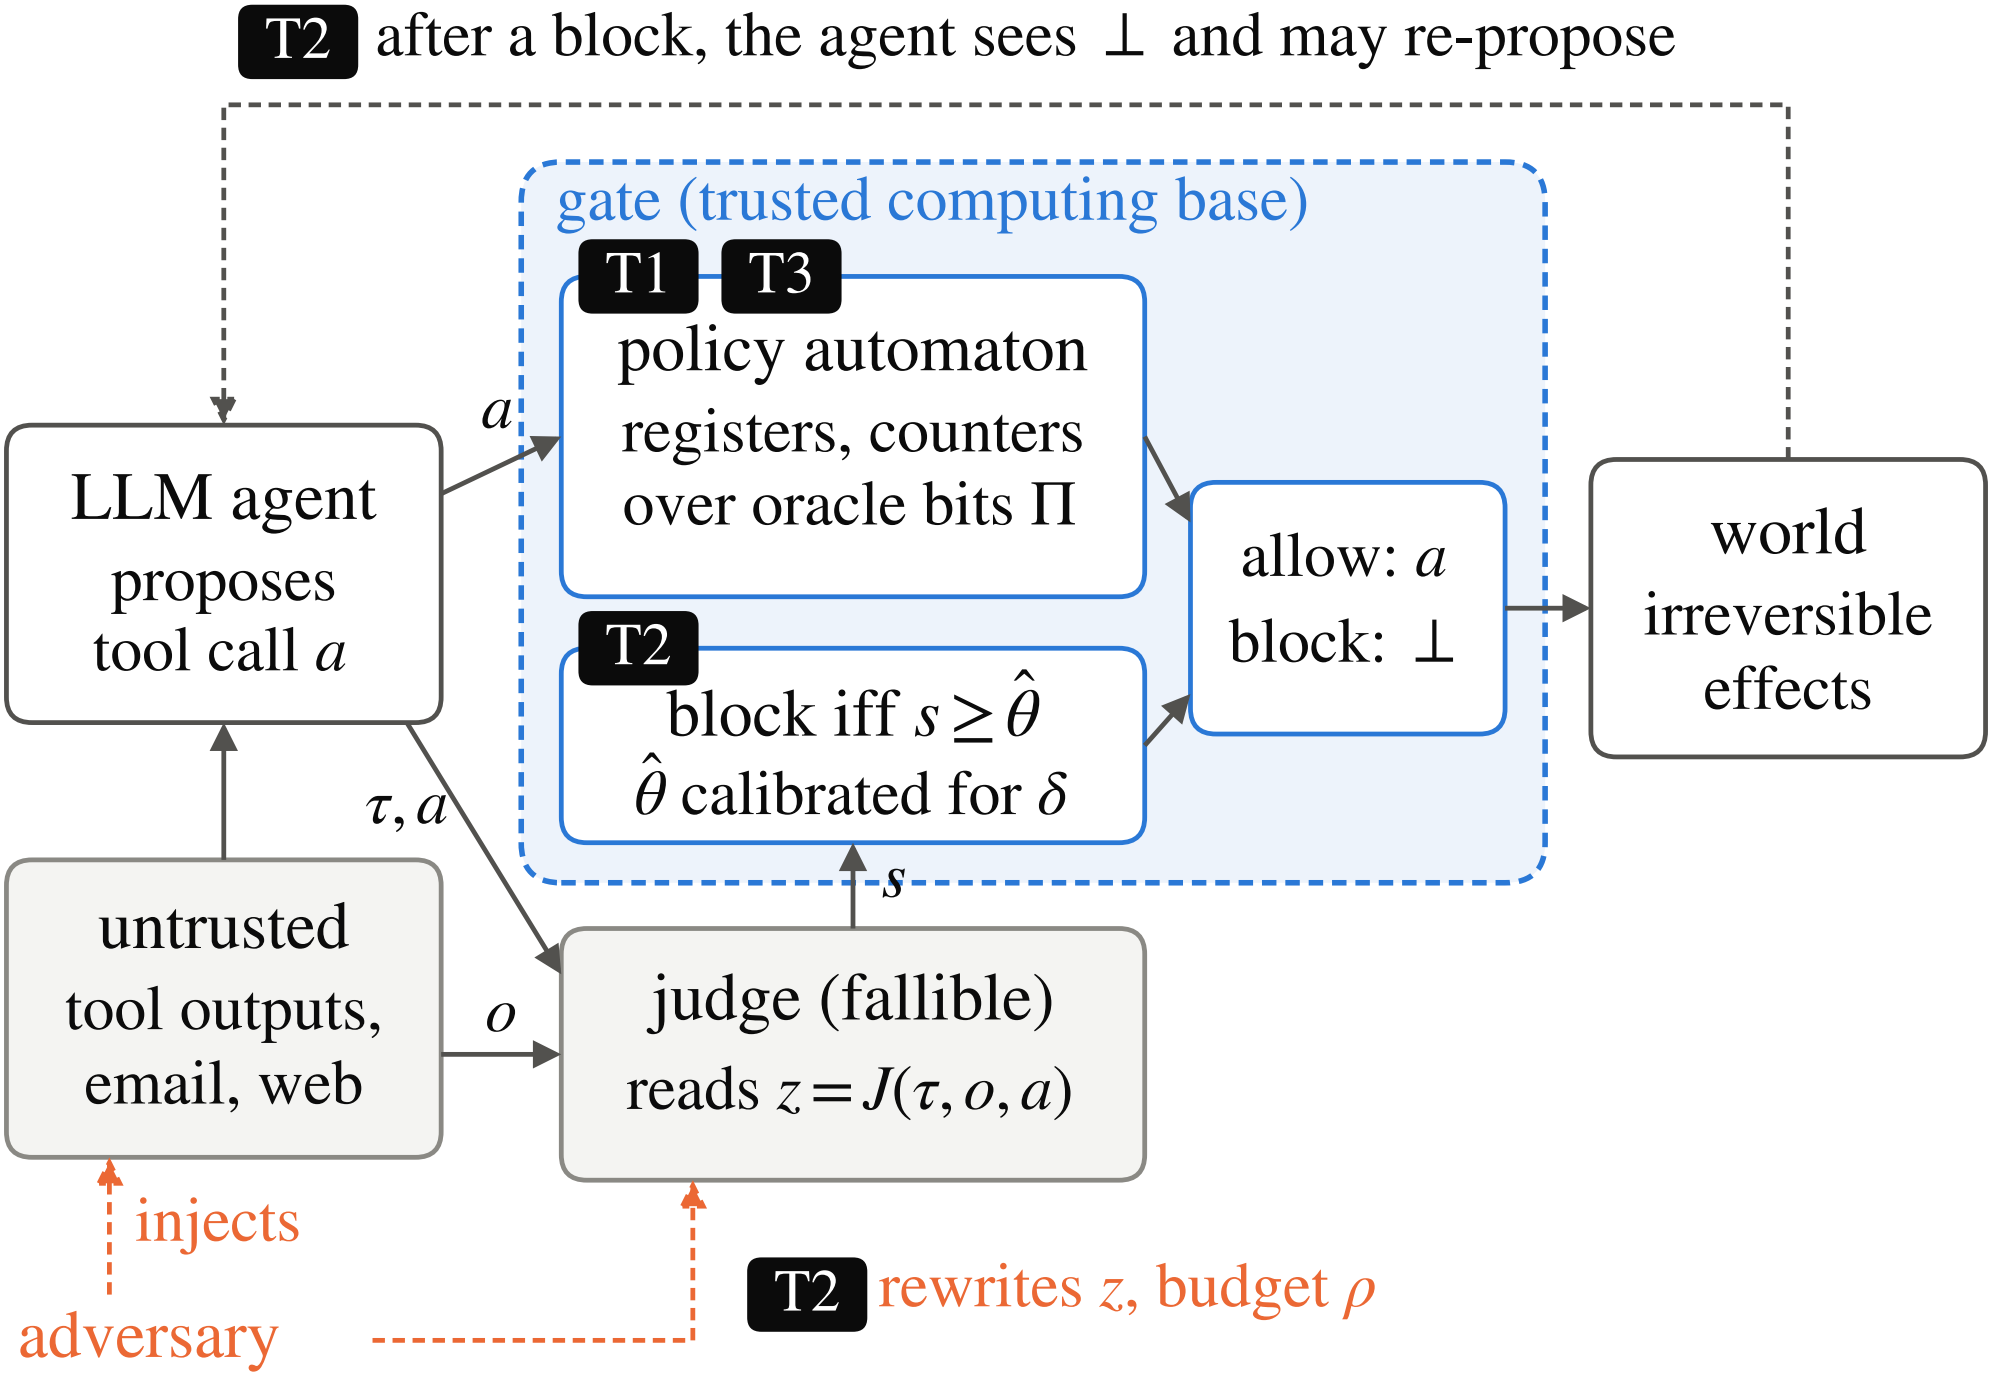
T2 after a block, the agent sees ⊥ and may re-propose
gate (trusted computing base)
T1
T3
policy automaton
registers, counters
over oracle bits Π
T2
block iff $s \geq \hat{\theta}$
$\hat{\theta}$ calibrated for δ
allow: a
block: ⊥
world
irreversible
effects
LLM agent
proposes
tool call a
a
τ, a
s
untrusted
tool outputs,
email, web
o
judge (fallible)
reads $z = J(\tau, o, a)$
injects
adversary
T2 rewrites z, budget ρ